




44. Sitzung der offenen Community of Practice



Agenda



- Update (allgemein & persönlich)
- Input: „Gemma“
- Diskussion

Update (allgemein)



Hintergrund:


- Rechts abgebildet ist auszugsweise das aktuelle LMArena-Leaderboard für die Kategorie „German“ mit Blick auf die besten open-weights-Modelle
- je 1 pro Modellfamilie, zzgl. kleine Modelle (<30B) sowie die Top-3 proprietären Modelle

Bemerkenswert:

- Anthropic's [Claude-opus-4-6-thinking](#) wird an der Spitze abgelöst durch [Gemini-3-pro](#)
- Bestes Open-Weights-Modell bleibt [Deepseek-v4-pro](#) (1.6T param. total, 49B active)
- Effizientestes Open-Weights Modell bleibt [Qwen3.5-27b](#)
 - wobei aktuelle Open-Weights-Modellfamilien wie [Qwen3.6](#) oder die [Gemma-4](#) noch ausstehen
 - „Text (Overall)“ weist Gemma-4-Modelle aus (s.u.)
- Alle dargestellten Modelle sind bereits besser als das hervorragende [Gemma-3-12b-it](#) 🤖

Arena Score German
selected models on lmarena.ai as of Jun 10, 2025

gemma-3-12b-it (Google · Gemma)
gemma-3-27b-it (Google · Gemma)
command-r-plus-0.1 (Cohere · CC-BY-NC-4.0)
llama-4-maverick-17b-128e-instruct (Meta · Llama 4)
gpt-oss-120b (OpenAI · Apache 2.0)
nvidia-nemotron-3-super-120b-a12b (Nvidia · NVIDIA Open Model)
llama-4-scout-17b-16e-instruct (Meta · Llama)
qwq-32b (Alibaba · Apache 2.0)
mistral-small-3.1-24b-instruct-2503 (Mistral · Apache 2.0)
olmo-3.1-32b-instruct (AI2 · Apache 2.0)
athene-70b-0725 (NexusFlow · CC-BY-NC-4.0)
qwen-max-0919 (Alibaba · Qwen)
gpt-oss-20b (OpenAI · Apache 2.0)
gemma-2-27b-it (Google · Gemma license)
qwen2.5-72b-instruct (Alibaba · Qwen)
mistral-small-24b-instruct-2501 (Mistral · Apache 2.0)
phi-4 (Microsoft · MIT)
jamba-1.5-large (AI21 Labs · Jamba Open)
nemotron-4-340b-instruct (Nvidia · NVIDIA Open Model)
c4ai-aya-expense-32b (Cohere · CC-BY-NC-4.0)
qwen2-72b-instruct (Alibaba · Qianwen LICENSE)
mistral-8x22b-instruct-v0.1 (Mistral · Apache 2.0)
qwen1.5-110b-chat (Alibaba · Qianwen LICENSE)
qwen1.5-4b-chat (Alibaba · Qianwen LICENSE)



Update (persönlich)



- Comic „das Kollegium“ mit KI generiert
- Beitrag zu KI in der beruflichen Orientierung und Beratung eingereicht (auf Basis unserer Evaluation des E-KI-BB- Projekts, s.Abbildung)
- Beitrag „Kompetenzprofile und Tätigkeitsanforderungen in der Transformation der Automobilindustrie. Ein KI-basierter Analyseansatz“ akzeptiert zur Publikation in der Zeitschrift für Berufs- und Wirtschaftspädagogik (ZBW)



Gemma (language model)

Article

Talk

Read Edit View history Tools

From Wikipedia, the free encyclopedia

Gemma is a series of source-available [large language models](#) developed by [Google DeepMind](#). It is based on similar technologies as [Gemini](#). The first version was released in February 2024, followed by **Gemma 2** in June 2024, **Gemma 3** in March 2025, and the [free and open-source](#) **Gemma 4** in April 2026. Variants of Gemma have also been developed, such as the vision-language model **PaliGemma** and the model **MedGemma** for medical consultation topics.

History [edit]

In February 2024, Google debuted Gemma, a collection of [source-available](#) LLMs that serve as a lightweight version of Gemini. The initial release came in two sizes: neural networks with two and seven billion [parameters](#) respectively. Multiple publications viewed this as a response to competitors such as [Meta](#) releasing source code for their AI models, and a shift from Google's longstanding practice of keeping its AI source code private.^{[4][5][6][7]}

Gemma 2 was released on June 27, 2024,^[8] and Gemma 3 was released on March 12, 2025.^{[9][10]} On April 2, 2026, Google released Gemma 4 under the free and open-source [Apache 2.0](#) license.^{[2][11]}

Gemma



Developer	Google DeepMind
Initial release	February 21, 2024; 2 years ago ^[1]
Stable release	Gemma 4 / April 2, 2026; 2 months ago ^[2]
Type	Large language model
License	Gemma 4: Apache 2.0 Gemma 3 and earlier: Source-available (Gemma Terms of Use) ^[3]
Website	deepmind.google/models/gemma/ 

Gemma: Open Models Based on Gemini Research and Technology

Gemma Team, Google DeepMind¹

This work introduces Gemma, a family of lightweight, state-of-the-art open models built from the research and technology used to create Gemini models. Gemma models demonstrate strong performance across academic benchmarks for language understanding, reasoning, and safety. We release two sizes of models (2 billion and 7 billion parameters), and provide both pretrained and fine-tuned checkpoints. Gemma outperforms similarly sized open models on 11 out of 18 text-based tasks, and we present comprehensive evaluations of safety and responsibility aspects of the models, alongside a detailed description of model development. We believe the responsible release of LLMs is critical for improving the safety of frontier models, and for enabling the next wave of LLM innovations.

Die erste Gemma-Generation (2B & 7B) erhielt am 5.4. ein Update auf v1.1, sprach allerdings noch kaum Deutsch

Google DeepMind

2024-06-27

Gemma 2: Improving Open Language Models at a Practical Size

Gemma Team, Google DeepMind¹

In this work, we introduce Gemma 2, a new addition to the Gemma family of lightweight, state-of-the-art open models, ranging in scale from 2 billion to 27 billion parameters. In this new version, we apply several known technical modifications to the Transformer architecture, such as interleaving local-global attentions (Beltagy et al., 2020a) and group-query attention (Ainslie et al., 2023). We also train the 2B and 9B models with knowledge distillation (Hinton et al., 2015) instead of next token prediction. The resulting models deliver the best performance for their size, and even offer competitive alternatives to models that are 2-3× bigger. We release all our models to the community.

Die zweite Gemma-Generation (9B & 27B, seit 31.07.24 auch 2B) führte auch zum grandiosen Gemma-2-9b-it-SimPo!

Gemma 3 Technical Report

Gemma Team, Google DeepMind¹

We introduce Gemma 3, a multimodal addition to the Gemma family of lightweight open models, ranging in scale from 1 to 27 billion parameters. This version introduces vision understanding abilities, a wider coverage of languages and longer context – at least 128K tokens. We also change the architecture of the model to reduce the KV-cache memory that tends to explode with long context. This is achieved by increasing the ratio of local to global attention layers, and keeping the span on local attention short. The Gemma 3 models are trained with distillation and achieve superior performance to Gemma 2 for both pre-trained and instruction finetuned versions. In particular, our novel post-training recipe significantly improves the math, chat, instruction-following and multilingual abilities, making Gemma3-4B-IT competitive with Gemma2-27B-IT and Gemma3-27B-IT comparable to Gemini-1.5-Pro across benchmarks. We release all our models to the community.

Die dritte Gemma-Generation (1B, 4B, 12B und 27B, später 270M) wurde ab 26.06. ergänzt um Gemma 3n (E2B & E4B)

Gemma 4 models are multimodal, handling text and image input (with audio supported on E2B, E4B, and 12B models) and generating text output.

Core Capabilities

Gemma 4 models handle a broad range of tasks across text, vision, and audio. Key capabilities include:


- **Thinking** – Built-in reasoning mode that lets the model think step-by-step before answering.
- **Long Context** – Context windows of up to 128K tokens (E2B/E4B) and 256K tokens (12B/26B A4B/31B).
- **Image Understanding** – Object detection, Document/PDF parsing, screen and UI understanding, chart comprehension, OCR (including multilingual), handwriting recognition, and pointing. Images can be processed at variable aspect ratios and resolutions.
- **Video Understanding** – Analyze video by processing sequences of frames.
- **Interleaved Multimodal Input** – Freely mix text and images in any order within a single prompt.
- **Function Calling** – Native support for structured tool use, enabling agentic workflows.
- **Coding** – Code generation, completion, and correction.
- **Multilingual** – Out-of-the-box support for 35+ languages, pre-trained on 140+ languages.
- **Audio** (E2B, E4B, and 12B Unified only) – Automatic speech recognition (ASR) and speech-to-translated-text translation across multiple languages.

Die vierte Gemma-Generation (E2B, E4B, 31B, 26 A4B; seit 03.06. auch 12B) erstmals unter Apache-2.0-Lizenz

Gemma 4 12B



- A multimodal encoder-free architecture: Bypassing heavy multi-stage vision and audio encoders entirely, multimodal data is fed straight into the LLM backbone, reducing multimodal latency.
- Our first medium-sized model with audio input: In the Gemma family, audio inputs were restricted to small, lightweight edge architectures (e.g. E4B). Gemma 4 12B is the first medium-sized model capable of natively ingesting audio.
- Developer-friendly size: Small enough to run locally on dedicated GPU laptops with 16GB VRAM or unified memory. To maximize local inference speeds, we are additionally releasing a dedicated multi-token prediction (MTP) model.
- New MacOS desktop experience: For the first time, we are releasing downloadable macOS desktop applications, letting developers experience fully local spoken and visual interaction directly on consumer-grade devices.



Quelle: <https://developers.googleblog.com/gemma-4-12b-the-developer-guide/>

Gemma 4 – Dense Models



Dense Models

Property	E2B	E4B	12B Unified	31B Dense
Total Parameters	2.3B effective (5.1B with embeddings)	4.5B effective (8B with embeddings)	11.95B	30.7B
Layers	35	42	48	60
Sliding Window	512 tokens	512 tokens	1024 tokens	1024 tokens
Context Length	128K tokens	128K tokens	256K tokens	256K tokens
Vocabulary Size	262K	262K	262K	262K
Supported Modalities	Text, Image, Audio	Text, Image, Audio	Text, Image, Audio	Text, Image
Vision Encoder Parameters	~150M	~150M	-	~550M
Audio Encoder Parameters	~300M	~300M	-	No Audio


The "E" in E2B and E4B stands for "effective" parameters. The smaller models incorporate Per-Layer Embeddings (PLE) to maximize parameter efficiency in on-device deployments. Rather than adding more layers or parameters to the model, PLE gives each decoder layer its own small embedding for every token. These embedding tables are large but are only used for quick lookups, which is why the effective parameter count is much smaller than the total.

Gemma 4 Benchmarks



Benchmark			Gemma 4 31B IT Thinking	Gemma 4 26B A4B IT Thinking	Gemma 4 E4B IT Thinking	Gemma 4 E2B IT Thinking	Gemma 3 27B IT
Arena AI (text) As of 4/2/26			1452	1441	-	-	1365
MMMLU Multilingual Q&A	No tools		85.2%	82.6%	69.4%	60.0%	67.6%
MMMU Pro Multimodal reasoning			76.9%	73.8%	52.6%	44.2%	49.7%
AIME 2026 Mathematics	No tools		89.2%	88.3%	42.5%	37.5%	20.8%
LiveCodeBench v6 Competitive coding problems			80.0%	77.1%	52.0%	44.0%	29.1%
GPQA Diamond Scientific knowledge	No tools		84.3%	82.3%	58.6%	43.4%	42.4%
τ2-bench Agentic tool use	Retail		86.4%	85.5%	57.5%	29.4%	6.6%

Gemma 4 12B - Benchmarks



Gemma 1-4 der Text-Arena



- Mit Blick auf die Kategorie „Text“ liegen die evaluierten Gemma-4-Modelle deutlich vor Gemma 3
- Aktuell sind noch nicht alle Modelle evaluiert (und noch keines für die Unterkategorie „German“)

Rank ↕	Rank Spread ↕	Model ↕	Score ↓	Votes ↕	Price \$/M ↕	Context ↕
41	30 → 58	gemma-4-31b Google · Apache 2.0	1451 ±8	5.881	\$0.14 / \$0.40	262.1K
59	42 → 78	gemma-4-26b-a4b Google · Apache 2.0	1438 ±8	5.804	N/A	N/A
155	147 → 163	gemma-3-27b-it Google · Gemma	1366 ±4	47.545	\$0.08 / \$0.16	131.1K
185	164 → 201	gemma-3-12b-it Google · Gemma	1342 ±10	3.829	\$0.05 / \$0.15	131.1K
214	198 → 229	gemma-3n-e4b-it Google · Gemma	1318 ±5	22.601	\$0.06 / \$0.12	32.8K
236	216 → 246	gemma-3-4b-it Google · Gemma	1303 ±9	4.171	\$0.05 / \$0.10	131.1K
243	238 → 251	gemma-2-27b-it Google · Gemma license	1289 ±3	75.754	\$0.65 / \$0.65	8.2K
252	240 → 260	gemma-2-9b-it-simpo Princeton · MIT	1280 ±7	10.072	\$0.03 / \$0.09	8.2K
262	256 → 268	gemma-2-9b-it Google · Gemma license	1266 ±4	54.611	\$0.03 / \$0.09	8.2K
296	292 → 301	gemma-2-2b-it Google · Gemma license	1200 ±4	46.616	N/A	N/A
308	300 → 317	gemma-1.1-7b-it Google · Gemma license	1181 ±6	23.893	\$0.03 / \$0.09	8.2K
335	321 → 344	gemma-7b-it Google · Gemma license	1136 ±9	8.925	\$0.05 / \$0.08	8.2K
344	337 → 349	gemma-1.1-2b-it Google · Gemma license	1115 ±8	10.854	N/A	N/A
350	345 → 353	gemma-2b-it Google · Gemma license	1092 ±12	4.780	\$0.10 / \$0.10	N/A

Gemma 1-3 und die deutsche Sprache



Zum Vergleich:
Mixtral-8x22b-
instruct-v0.1
lag bei 1217

Rank ↕	Rank Spread ↕	Model ↕	Score ↓	Votes ↕	Price \$/M ↕ ↕	Context ↕
120	48 ↔ 171	 gemma-3-12b-it Google · Gemma	1371 ±43	166	\$0.05 / \$0.15	131.1K
121	78 ↔ 150	 gemma-3-27b-it Google · Gemma	1371 ±18	1.164	\$0.08 / \$0.16	131.1K
157	126 ↔ 199	 gemma-3n-e4b-it Google · Gemma	1316 ±22	692	\$0.06 / \$0.12	32.8K
185	131 ↔ 225	 gemma-3-4b-it Google · Gemma	1289 ±42	162	\$0.05 / \$0.10	131.1K
198	151 ↔ 230	 gemma-2-9b-it-simpo Princeton · MIT	1269 ±31	318	\$0.03 / \$0.09	8.2K
199	170 ↔ 222	 gemma-2-27b-it Google · Gemma license	1269 ±14	2.037	\$0.65 / \$0.65	8.2K
217	186 ↔ 232	 gemma-2-9b-it Google · Gemma license	1245 ±16	1.580	\$0.03 / \$0.09	8.2K
241	235 ↔ 254	 gemma-2-2b-it Google · Gemma license	1155 ±19	1.158	N/A	N/A
243	233 ↔ 257	 gemma-1.1-7b-it Google · Gemma license	1155 ±22	827	\$0.03 / \$0.09	8.2K
267	254 ↔ 270	 gemma-1.1-2b-it Google · Gemma license	1048 ±35	365	N/A	N/A

- Fragen?
- Anregungen?
- Erfahrungen?